

LINUX CLUSTER MONITOR

- Two Years' Progress Review -

HPC & AI Team

June 2nd 2022

Chenyang Li, Mario Rodriguez

Reeann Zhang



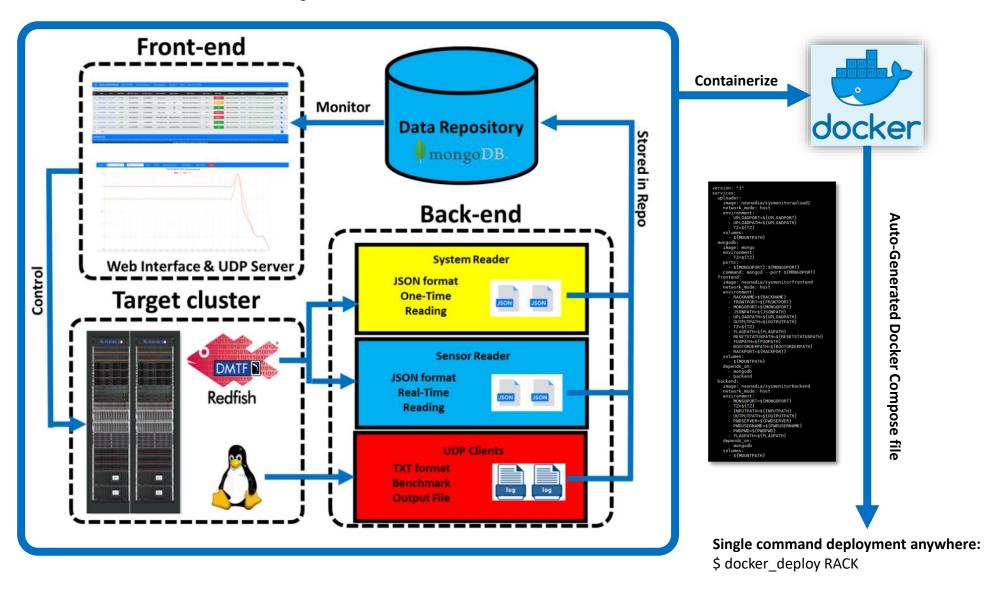
Outline



- ☐ System architecture
- ☐ Latest updates overview
- ☐ Web client
- ☐ Advanced features
- ☐ Sensor reading features
- ☐ Cluster report generator
- ☐ Quick deployment steps

System Architecture





Outline



- ☐ System architecture
- ☐ Latest updates overview
- ☐ Web client
- ☐ Advanced features
- ☐ Sensor reading features
- ☐ Cluster report generator
- ☐ Quick deployment steps

Latest Updates Overview



Front-end:

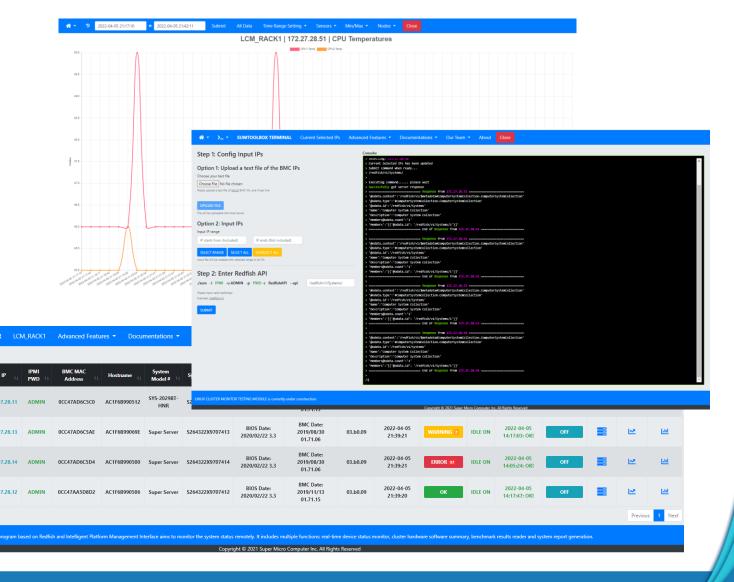
- 1. Asynchronous web page.
- 2. Real-time rack level administration tools.
- 3. More advanced features.
- 4. New cluster report.
- 5. User friendly UI design.

Back-end:

- 1. Multiple APIs: Redfish, IPMI and SUM.
- 2. Performance optimization.
- 3. Better benchmark support.
- 4. Archive module.

Deployment:

- 1. Automatic deployment.
- 2. Input file auto generating.
- 3. Health check and autoheal.



Outline



- ☐ System architecture
- ☐ Latest updates overview
- ☐ Web client
- ☐ Advanced features
- ☐ Sensor reading features
- ☐ Cluster report generator
- ☐ Quick deployment steps

Web Client: Rack View



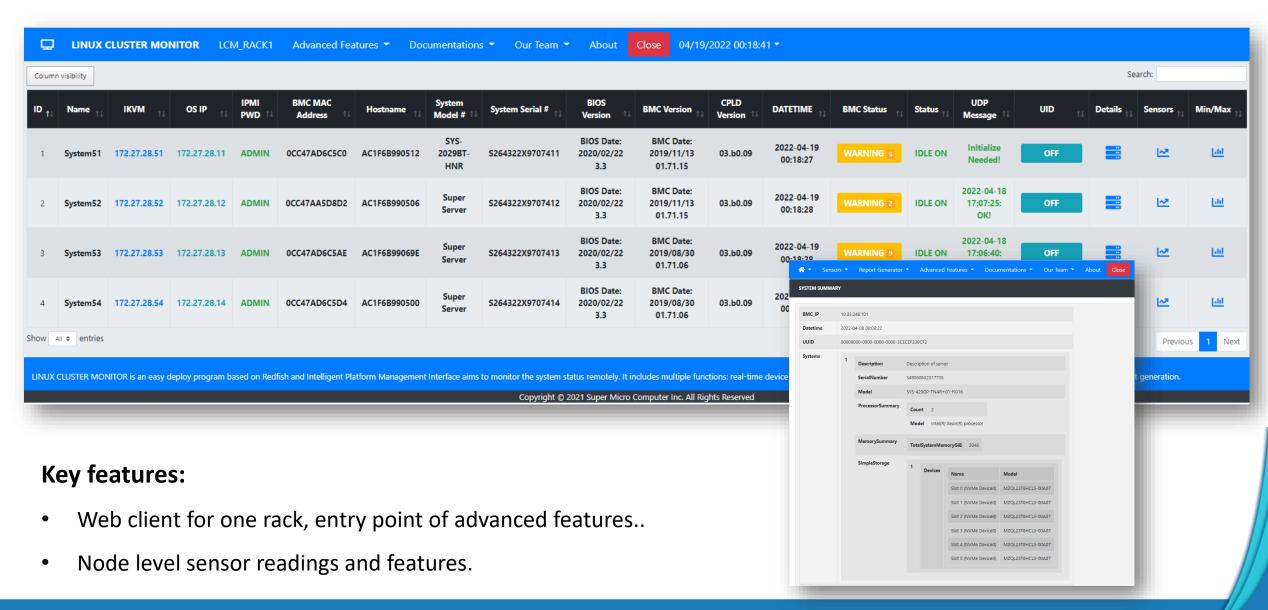
	mn visibility Sea																	
D 11	RACK NAME	NODES †1	START DATE	LAST DATE †1	TIME ELAPSED 11	IPMI RANGE	BMC EVENTS	POWER †1	TEMP †1	FAN †1	MIN/MAX _† ↓	IPMI AUTH 1	REPORT 11	SUM †1	IPMITOOL †1	SENSOR COUNT	⊥ UDP †↓	DETAIL
1	N1-rack1-mapping_SSG-6029P- E1CR24L	11	2021-08-27 11:56:52	2021-08-30 10:09:03	2 days, 22:12:11	192.168.8.237 - 192.168.13.30	ОК	₩.	<u>~</u>	₩.	<u>lad</u>	csv	ß	<u>&</u>	&	¥≡	&	==0 ==0 ==0
2	N2-rack1-mapping_SYS-1029P- WTRT	5	2021-08-27 11:57:03	2021-08-30 10:08:25	2 days, 22:11:22	192.168.8.205 - 192.168.13.22	ОК	<u>~</u>	<u>~</u>	₩.	<u>lad</u>	csv	ß	&	A	¥≡	&	- 17 - 12 - 13
3	N3-rack1-mapping_SYS-1019P- WTR	5	2021-08-27 11:56:57	2021-08-30 10:08:50	2 days, 22:11:53	192.168.9.2 - 192.168.14.9	ОК	<u>~</u>	<u>~</u>	<u>~</u>	<u>lad</u>	csv	ß	8	a	¥≡	&	10 11
4	N4-rack1-mapping_SYS-6029U- E1CRT	2	2021-08-27 11:57:08	2021-08-30 10:08:56	2 days, 22:11:48	192.168.9.59 - 192.168.9.253	ОК	₩.	<u>~</u>	₩.	<u>lad</u>	csv	ß	&	&	¥≣	&	10
5	N5-rack10-mapping-new_SYS- 2029BT-HNC0R	36	2021-08-26 16:22:13	2021-09-01 09:28:02	5 days, 17:05:49	192.168.0.24 - 192.168.15.248	2 Errors 34 Warnings	<u>~</u>	<u>~</u>	₩.	<u>lad</u>	csv	ß	&	&	¥≣	&	11
6	N6-rack10-mapping-new_SYS- 4029GP-TRT	2	2021-08-26 16:21:56	2021-09-01 09:28:36	5 days, 17:06:40	192.168.0.163 - 192.168.14.149	2 Warnings	₩.	<u>~</u>	₩.	<u>[ad</u>	csv	ß	<u>&</u>	&	¥≣	&	11
All ¢	entries																Previous	s 1 I

- Web client for everything, each row represent a single rack.
- Rack level sensor readings and features.



Web Client: Node View





Outline



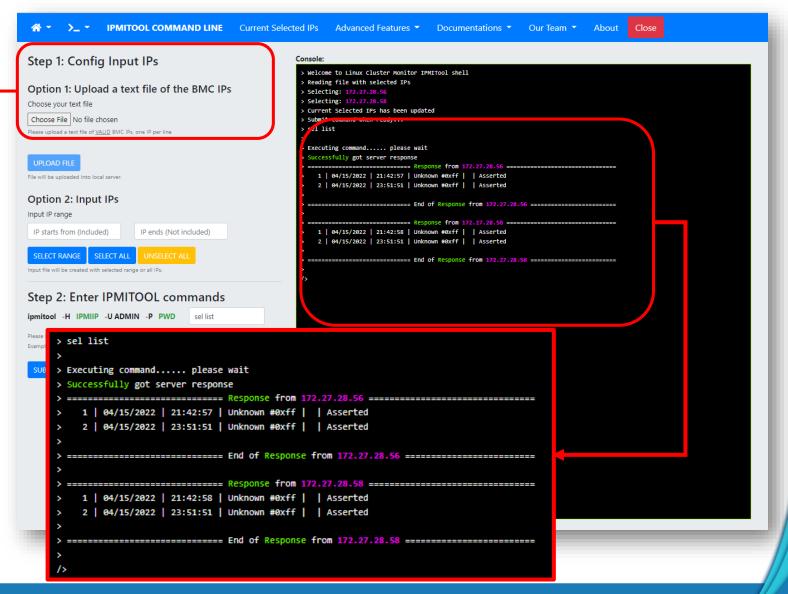
- ☐ System architecture
- ☐ Latest updates overview
- ☐ Web client
- Advanced features
- ☐ Sensor reading features
- ☐ Cluster report generator
- ☐ Quick deployment steps



IPMI COMMAND LINE



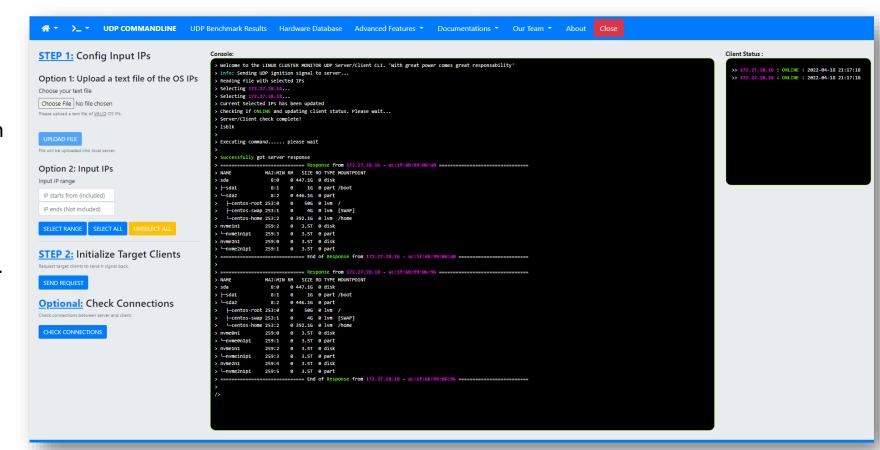
- Ease to execute ipmi commands
- A top down message is generated with the results from your nodes.
- Select your nodes with a text file, manually input range.
- User and Password must be correct on deployment of Linux Cluster Monitor.

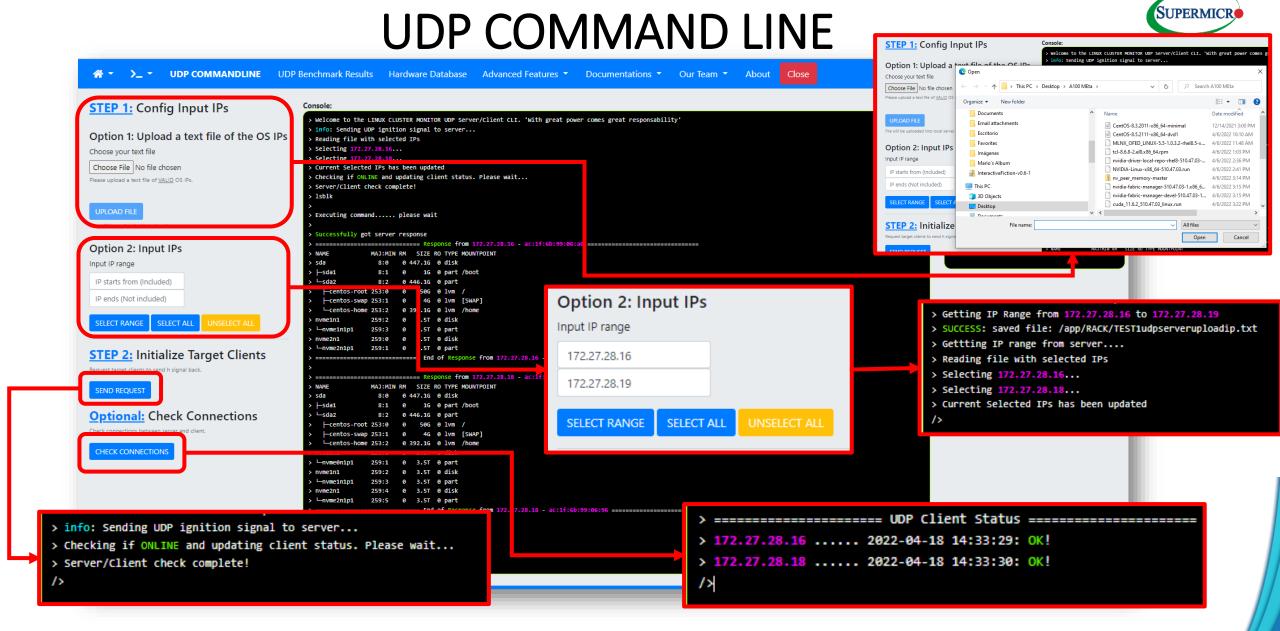




UDP COMMAND LINE

- Run bash commands through UDP.
- UDP server will initialize the clients in your cluster on page request.
- UDP clients must be installed and initialized on your cluster by the user.
- CLI gives constant feedback.
- Connection check between client/server.

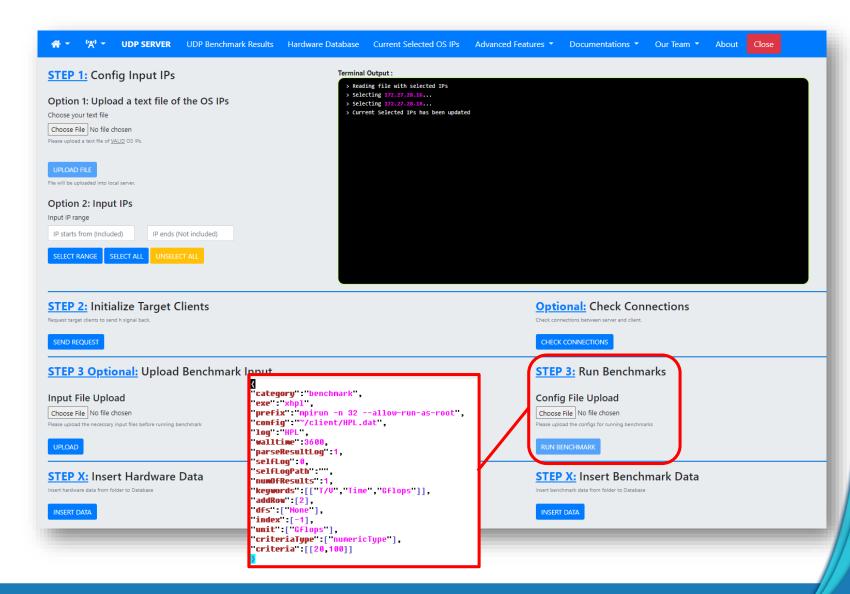






UDP SERVER CONTROLLER

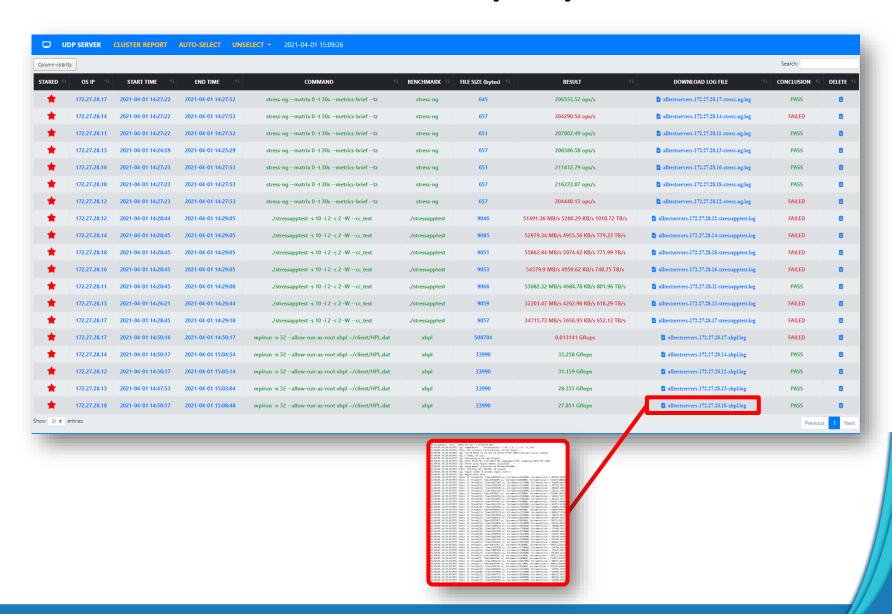
- All benchmarks supported.
- Config file needed for each benchmark.
- Benchmark results will be saved to database once done.
- Allows user to load hardware data from your cluster
- Allows user to load benchmark data into the database for parsing.
- Real-time terminal output to show the running status.



UDP Benchmark Results Display

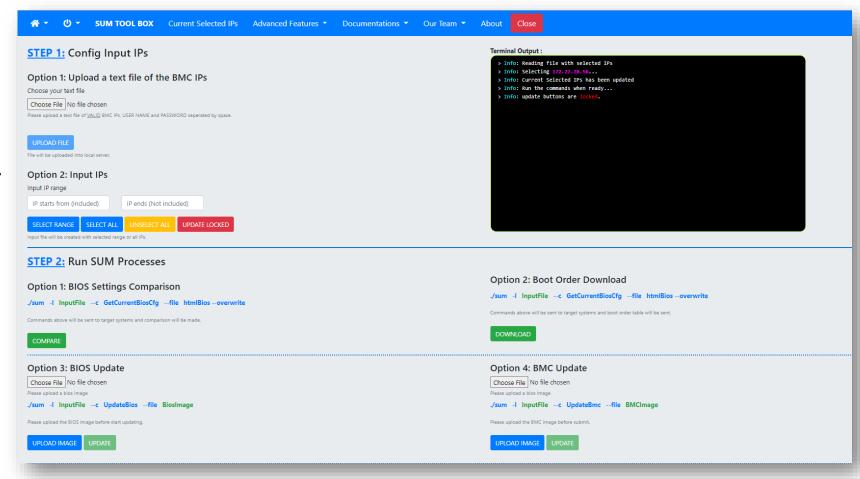


- Benchmark results are stored in database.
- Automatically fetch key results from output file.
- Automatically verify the performance.
- Start and end time of the benchmark.
- Benchmark type and runtime command.
- Cluster report generator.



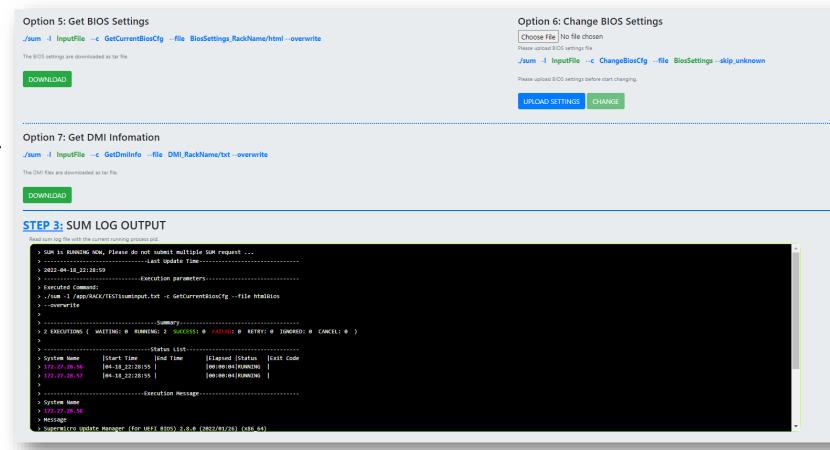


- Update BIOS.
- Update BMC firmware.
- Get BIOS settings for all nodes.
- Compare BIOS settings.
- Change BIOS settings
- Download boot order.
- Get DMI information.
- Show SUM running status.

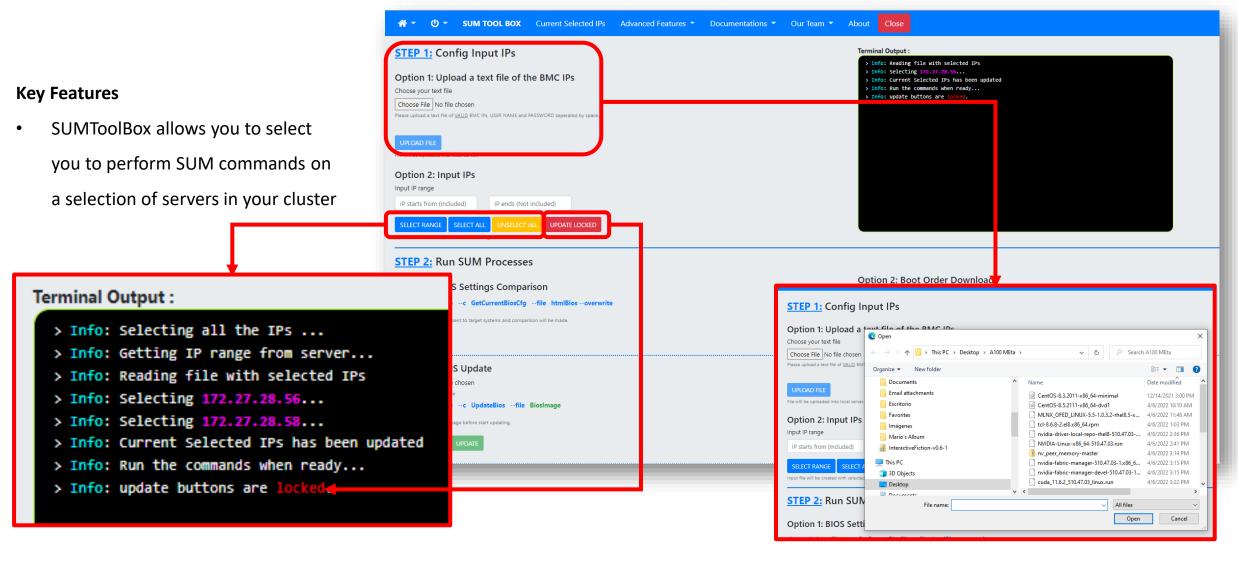




- Update BIOS.
- Update BMC firmware.
- Get BIOS settings for all nodes.
- Compare BIOS settings.
- Change BIOS settings
- Download boot order.
- Get DMI information.
- Show SUM LOG OUTPUT



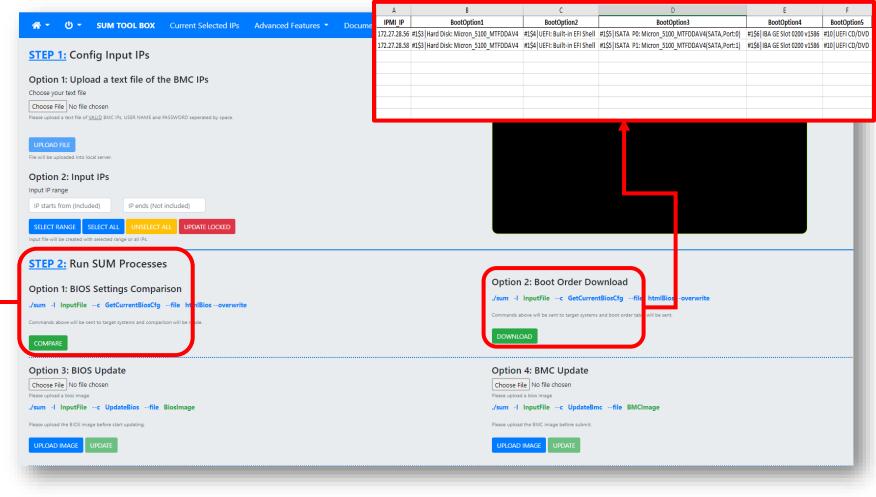






- BIOS Settings Comparison
- Boot Order Download

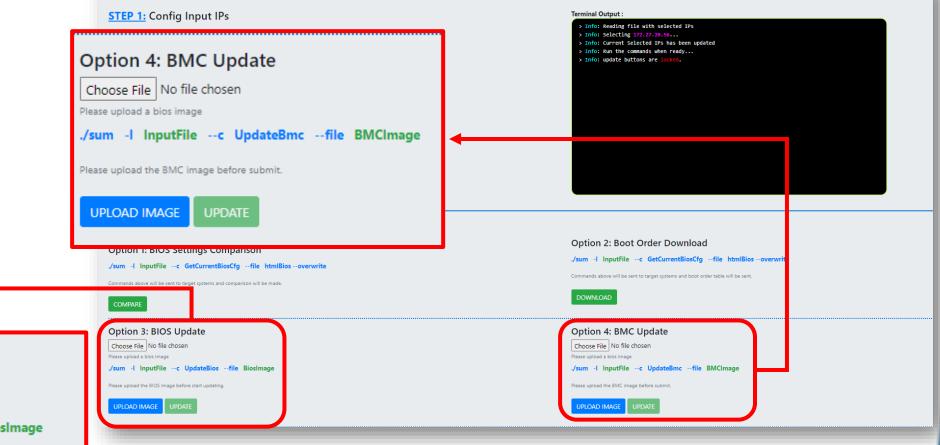




** U * SUM TOOL BOX Current Selected IPs Advanced Features * Documentations * Our Team * About Close



- BIOS update for multiple nodes
- BMC (IPMI firmware) update for multiple nodes
- Terminal to track the updating process.





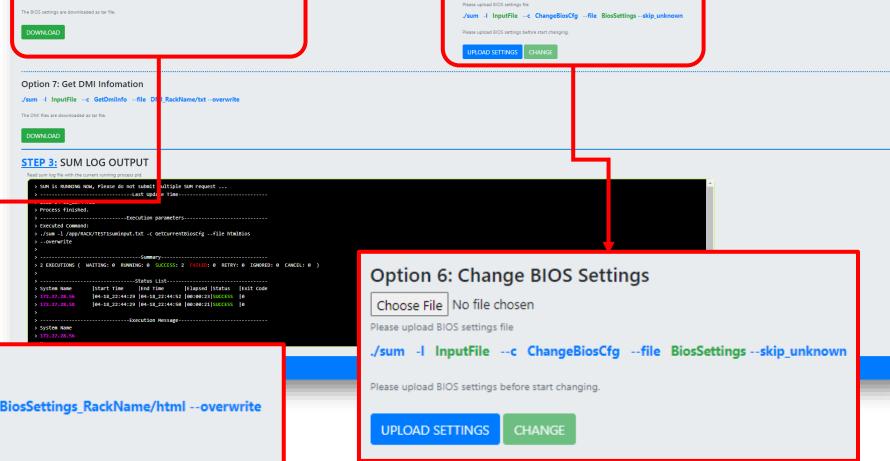
Option 5: Get BIOS Settings

/sum -I InputFile --c GetCurrentBiosCfg --file BiosSettings_RackName/html --overwrite



Key Features

- Get BIOS Settings and dump to html file.
- Change BIOS Settings by using html file.
- Terminal output to track.



Option 6: Change BIOS Settings

Choose File No file chosen

Option 5: Get BIOS Settings

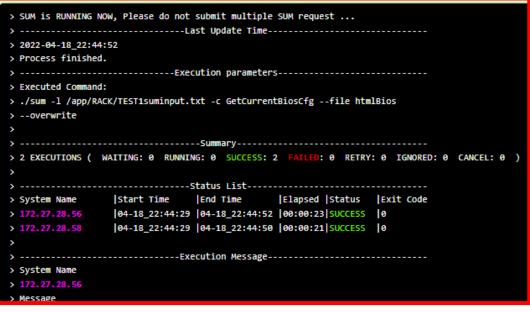
./sum -I InputFile --c GetCurrentBiosCfg --file BiosSettings_RackName/html --overwrite

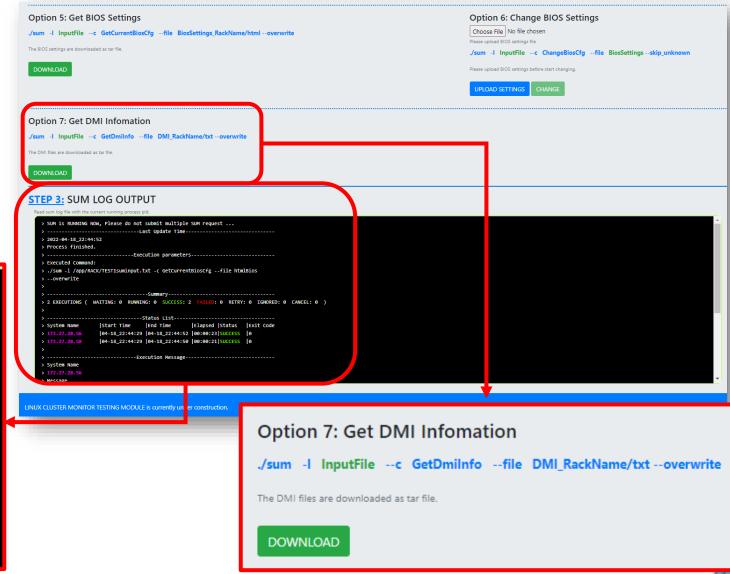
The BIOS settings are downloaded as tar file.

DOWNLOAD



- Get DMI Information
- SUM LOG OUTPUT







Key Features

/redfish/v1/Chassis

Executing command..... please wait

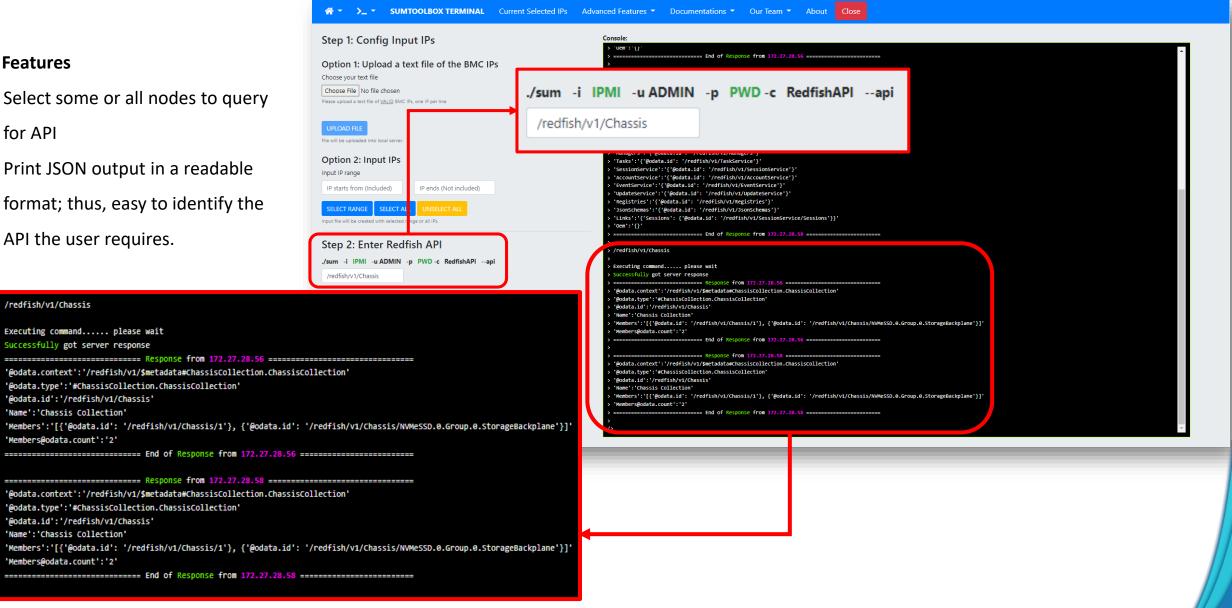
Successfully got server response

'@odata.id':'/redfish/v1/Chassis' 'Name':'Chassis Collection'

'Members@odata.count':'2'

'Members@odata.count':'2'

- Select some or all nodes to query for API
- Print JSON output in a readable format; thus, easy to identify the API the user requires.

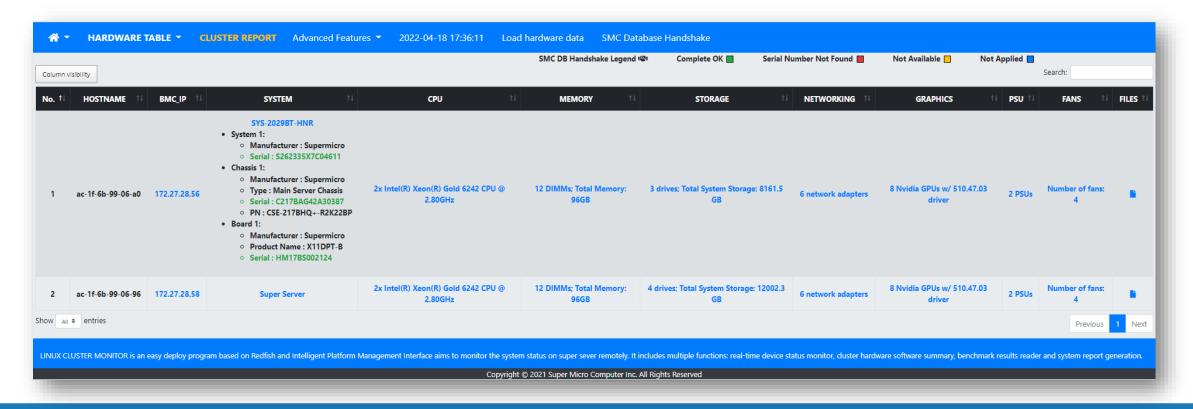


'@odata.id':'/redfish/v1/Chassis' 'Name':'Chassis Collection'

Rack Hardware viewer



- Parse through your hardware log for you cluster
- Create detailed table to show all the components and serial numbers.
- Handshake your MongoDB serial numbers w/ those from the SMC DB. SMC DB is provided as a .txt file corresponding to each system.
- Add your hardware data to an auto generated PDF Report.



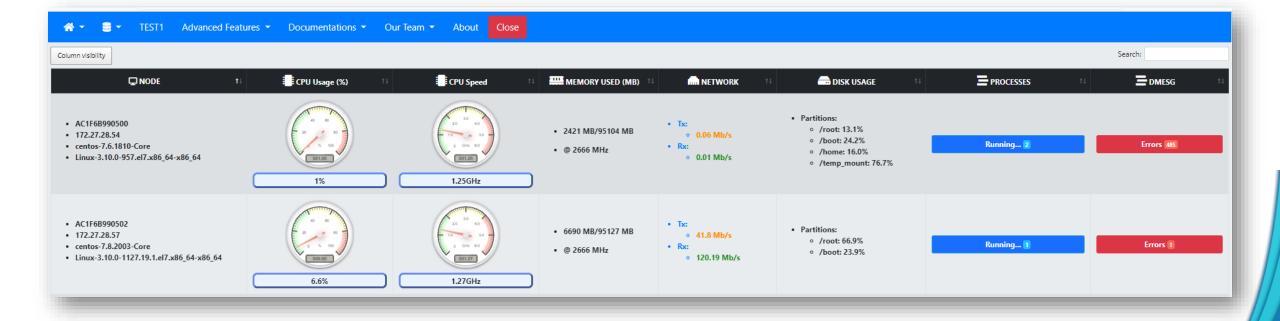
Cluster Telemetry



View server real-time metrics

- CPU Usage and Speed
- Partition usage
- Processes
- DMESG error capture
- Memory Usage

- Node kernel and Operating System
- Server Network Rx and Tx, with individualized NIC traffic mapped.



Outline

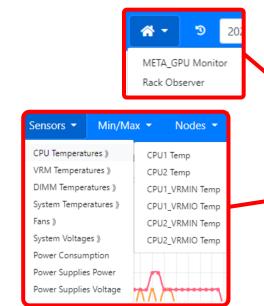


- ☐ System architecture
- ☐ Latest updates overview
- ☐ Web client
- ☐ Advanced features
- ☐ Sensor reading features
- ☐ Cluster report generator
- ☐ Quick deployment steps

Real-Time Sensor Readings

26





Key features:

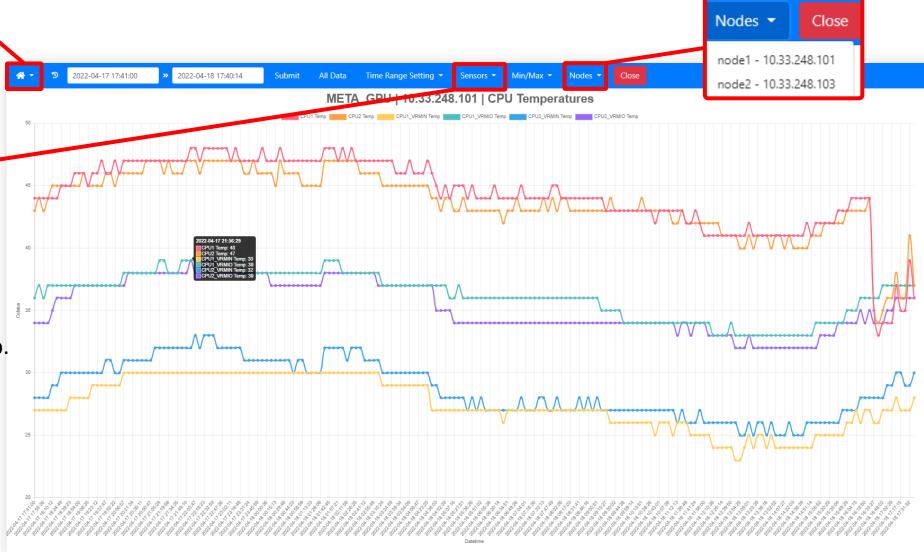
Line plots with timestamp.

Node selection.

Sensor selection.

New navigation button.

• Datetime backtracking.

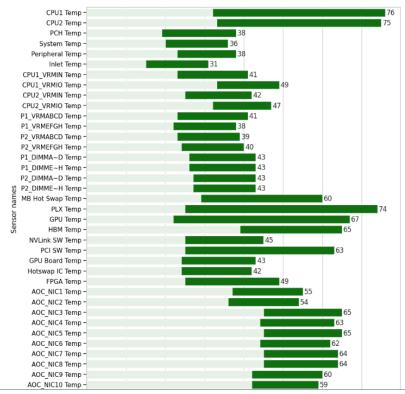


Supports power, temperature, voltage and fan speed.

Minimum/Maximum Sensor Readings







SENSOR NAME	MIN READING	MIN DATETIME	MAX READING	MAX DATETIME	AVG READING	NUM COUNT	NA COUNT
CPU1 Temp	32	2022-04-12 17:54:28	76	2022-04-11 22:22:39	46.2115	14653	164
CPU2 Temp	33	2022-04-13 16:33:13	75	2022-04-11 22:24:44	47.0535	14653	164
PCH Temp	19	2022-04-18 12:37:45	38	2022-04-08 00:27:23	25.716	14653	164
System Temp	20	2022-04-12 11:22:09	36	2022-04-08 00:08:38	25.569	14653	164
Peripheral Temp	23	2022-04-18 12:39:50	38	2022-04-08 00:07:35	29.2717	14653	164
Inlet Temp	15	2022-04-12 14:29:59	31	2022-04-08 00:05:29	20.8381	14653	164
CPU1_VRMIN Temp	23	2022-04-18 12:45:03	41	2022-04-11 19:43:04	30.3646	14653	164
CPU1_VRMIO Temp	33	2022-04-17 13:46:19	49	2022-04-08 00:03:24	38.9616	14653	164
CPU2_VRMIN Temp	25	2022-04-18 12:44:00	42	2022-04-11 21:27:21	31.7611	14653	164
CPU2_VRMIO Temp	32	2022-04-18 12:40:53	47	2022-04-08 00:07:35	37.9329	14653	164
P1_VRMABCD Temp	23	2022-04-18 12:41:55	41	2022-04-08 00:09:40	29.9841	14653	164
P1_VRMEFGH Temp	22	2022-04-18 12:41:55	38	2022-04-08 00:03:24	28.8001	14653	164
P2_VRMABCD Temp	23	2022-04-18 12:44:00	39	2022-04-08 00:03:24	29.8514	14653	164
P2_VRMEFGH Temp	24	2022-04-18 12:44:00	40	2022-04-08 00:24:15	30.0016	14653	164
P1_DIMMA~D Temp	26	2022-04-18 12:08:33	43	2022-04-08 00:06:33	32.5856	14653	164
P1_DIMME~H Temp	26	2022-04-18 12:50:15	43	2022-04-08 00:07:35	33.251	14653	164
P2_DIMMA~D Temp	27	2022-04-18 12:57:33	43	2022-04-08 00:03:24	34.0354	14653	164
P2_DIMME~H Temp	27	2022-04-13 14:22:51	43	2022-04-08 00:16:58	32.8862	14653	164
MB Hot Swap Temp	29	2022-04-18 12:49:12	60	2022-04-11 22:12:13	37.4129	14653	164
PLX Temp	25	2022-04-12 19:14:29	74	2022-04-12 17:51:20	57.1058	14653	164

- Bar plots with maximum and minimum readings.
- Table with numerical data and timestamp.
- Quick jump to the Real-Time sensor page with peak reading time period.
- Supports power, temperature, voltage and fan speed.
- LCM elapsed time and number recorded readings.

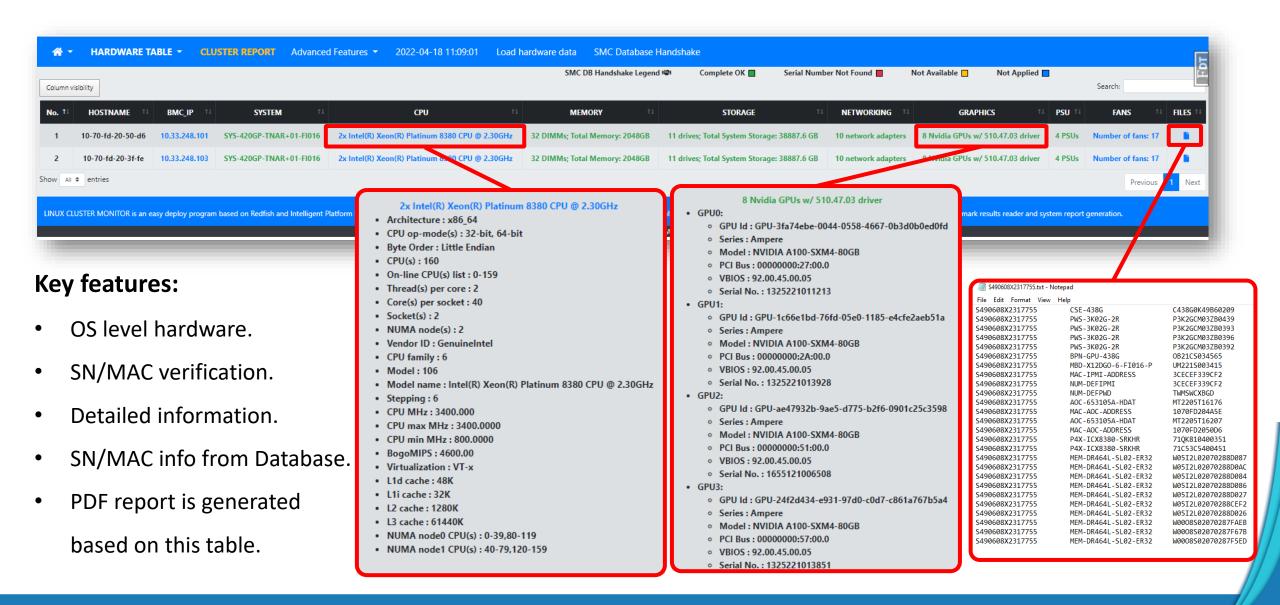
Outline



- ☐ System architecture
- ☐ Latest updates overview
- ☐ Web client
- ☐ Advanced features
- ☐ Sensor reading features
- ☐ Cluster report generator
- ☐ Quick deployment steps

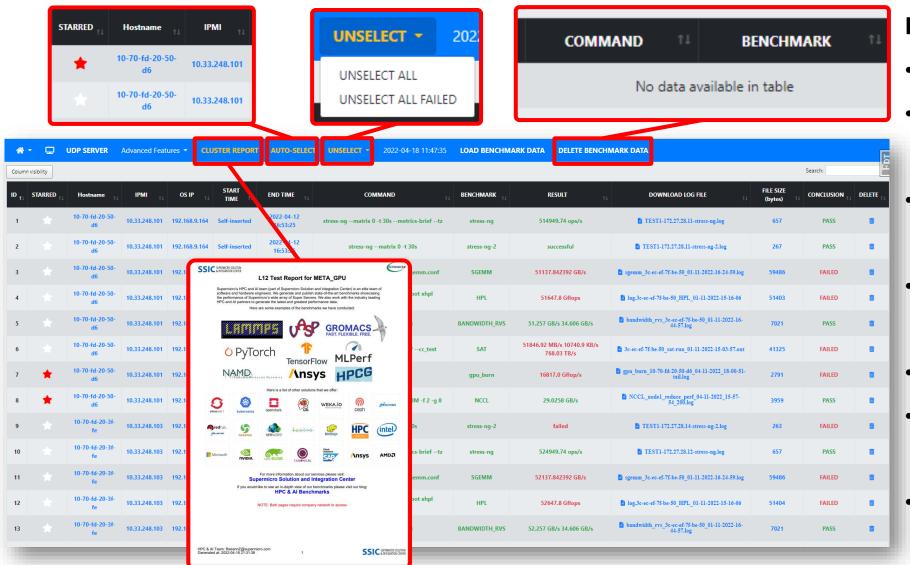
OS Level Hardware Information





Benchmark Results Table



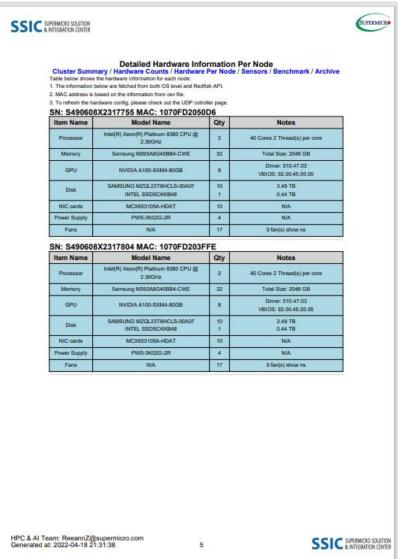


- Datetime of benchmark result.
- Benchmark conclusion: Pass or Failed.
- Benchmark running commands.
- Hostname and IPMI IP of benchmark node.
- Download log files.
- Select proper results for the report.
 - Results management.

Auto-Generated L12 Testing Report



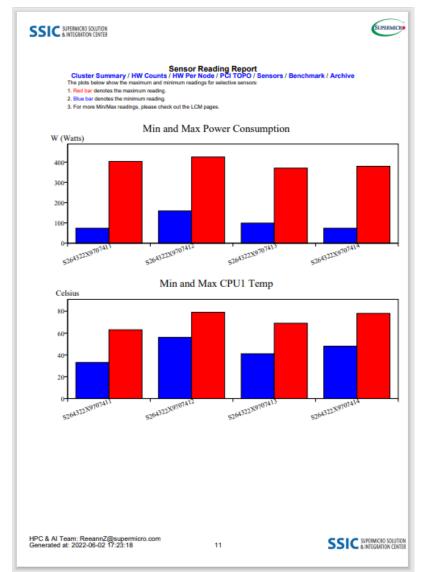


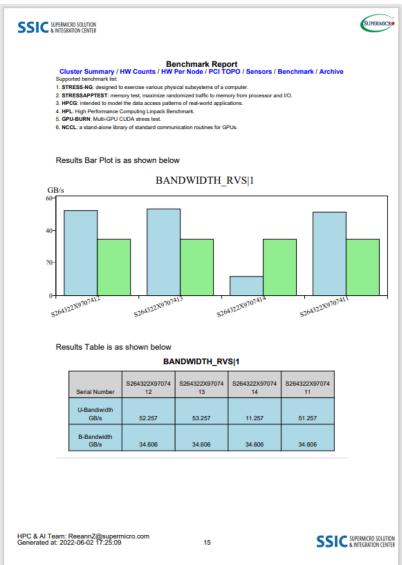


- Hardware information per node.
- Model name and quantity.
- Serial numbers and MAC addresses.
- A list of benchmarks we supported.
- Title page with introduction.
- Links to team's website.

Auto-Generated L12 Testing Report



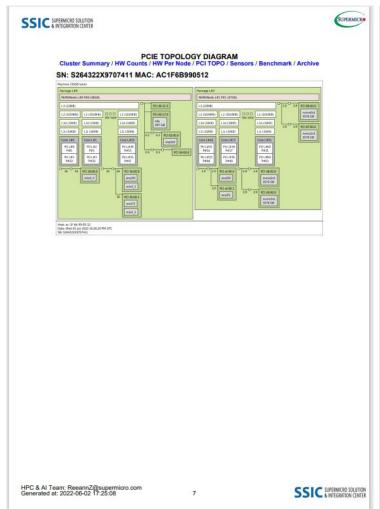


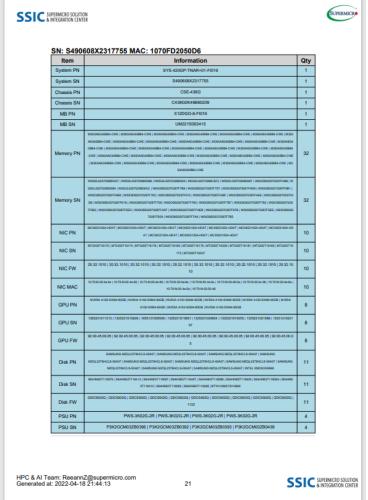


- Selected Min/Max sensor readings.
- Selected benchmark results.
- Each node is identified by serial number.
- Benchmark results have both table and charts.
- Currently support 6 kinds of benchmarks.
- Scalable for future benchmarks.

Auto-Generated L12 Testing Report







- Topology of all PCIE buses.
- Each table represent one node.
- Serial number and Model name for every component.
- Mac address for some of the components.
- FW for GPU, NIC and DISK.
- All the information is obtained from OS level or SMCIPMITool.

Outline



- ☐ System architecture
- ☐ Latest updates overview
- ☐ Web client
- ☐ Advanced features
- ☐ Sensor reading features
- ☐ Cluster report generator
- ☐ Quick deployment steps

Quick Deployment Steps



1. Install Docker, Docker-compose and ncat:

I. Docker: https://docs.docker.com/engine/install/centos/

II. Docker-compose: https://docs.docker.com/compose/install/

III. Ncat: \$ yum install nmap-ncat.x86_64

IV. Build necessary images, for the source code please visit our git server:

http://172.27.21.102:3000/ChenyangL/LinuxClusterMonitor

2. Bash shell script for deployment:

Following the link below and download the latest starting script: "docker_deploy.sh"

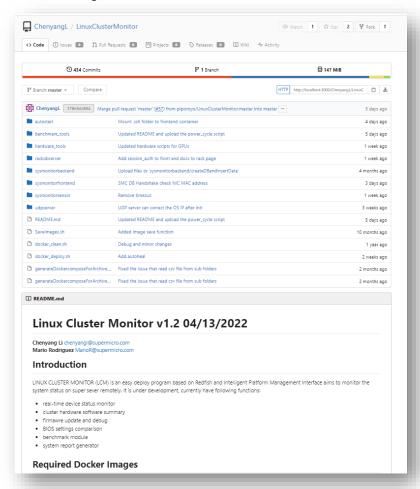
http://172.27.21.102:3000/ChenyangL/LinuxClusterMonitor

3. Create input files and start deployment:

- I. Create a folder for input files.
- II. Copy ".csv" files into folders.
- III. Create an "auto.env" file, visit our git server for more details.
- IV. Run: \$./docker_deploy.sh FOLDERNAME

4. Healthcheck and auto-heal:

- I. Healthcheck functions has been implemented to frontend pages and sensor reading module to monitoring the healthy status.
- II. If container is not healthy, it will be restarted. (auto-heal).
- III. To disable the auto-heal, simply stop the 'autoheal' container.





DISCLAIMER

Super Micro Computer, Inc. may make changes to specifications and product descriptions at any time, without notice. The information presented in this document is for informational purposes only and may contain technical inaccuracies, omissions and typographical errors. Any performance tests and ratings are measured using systems that reflect the approximate performance of Super Micro Computer, Inc. products as measured by those tests. Any differences in software or hardware configuration may affect actual performance, and Super Micro Computer, Inc. does not control the design or implementation of third party benchmarks or websites referenced in this document. The information contained herein is subject to change and may be rendered inaccurate for many reasons, including but not limited to any changes in product and/or roadmap, component and hardware revision changes, new model and/or product releases, software changes, firmware changes, or the like. Super Micro Computer, Inc. assumes no obligation to update or otherwise correct or revise this information.

SUPER MICRO COMPUTER, INC. MAKES NO REPRESENTATIONS OR WARRANTIES WITH RESPECT TO THE CONTENTS HEREOF AND ASSUMES NO RESPONSIBILITY FOR ANY INACCURACIES, ERRORS OR OMISSIONS THAT MAY APPEAR IN THIS INFORMATION.

SUPER MICRO COMPUTER, INC. SPECIFICALLY DISCLAIMS ANY IMPLIED WARRANTIES OF MERCHANTABILITY OR FITNESS FOR ANY PARTICULAR PURPOSE. IN NO EVENT WILL SUPER MICRO COMPUTER, INC. BE LIABLE TO ANY PERSON FOR ANY DIRECT, INDIRECT, SPECIAL OR OTHER CONSEQUENTIAL DAMAGES ARISING FROM THE USE OF ANY INFORMATION CONTAINED HEREIN, EVEN IF SUPER MICRO COMPUTER, Inc. IS EXPRESSLY ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

ATTRIBUTION

© 2022 Super Micro Computer, Inc. All rights reserved.

Thank You



www.supermicro.com